

Untitled

Beate, Marcus, Jonas, Wilmer

2024-04-11

First load the following libraries

```
library(tidyverse)
library(tidymodels)
library(dplyr)
library(rsample)
library(ranger)
```

The data is fetched from “<https://shorturl.at/crAK3>”. This dataset contains files for four wine styles: red, white, rosé, and sparkling. Each of these has eight columns, including the “NumberOfRatings” column indicating the count of people who rated the wine.

Each dataset is augmented with a “Type” column indicating the wine type, and unnecessary columns like “Name”, “Winery”, and “Region” are removed for analysis.

```
red <- read.csv("Red.csv")
white <- read.csv("White.csv")
rose <- read.csv("Rose.csv")
sparkling <- read.csv("Sparkling.csv")

red$Type <- "Red"
white$Type <- "White"
rose$Type <- "Rose"
sparkling$Type <- "Sparkling"

data <- rbind(red, white, rose, sparkling) %>% select(-Name, -Winery, -Region)
data$Country <- as.factor(data$Country)
data$Year <- as.factor(data$Year)
data$Type <- as.factor(data$Type)

data <- data %>% filter(Year != "N.V."); head(data)
```

```
##   Country Rating NumberOfRatings Price Year Type
## 1  France   4.2             100 95.00 2011  Red
## 2  France   4.3             100 15.50 2017  Red
## 3   Italy   3.9             100  7.45 2015  Red
## 4   Italy   3.5             100  8.72 2019  Red
## 5 Austria   3.9             100 29.15 2016  Red
## 6  France   3.7             100 19.90 2017  Red
```

Here, we split the combined dataset into training and testing sets. This step ensures that our models are trained on a portion of the data and evaluated on unseen data.

```
cell_split <- initial_split(data, prop = 3/70, strata = "Type")
```

```
train_data <- training(cell_split)
test_data <- testing(cell_split)
```

We construct a random forest regression model with 1000 trees and fit it to the training data. Additionally, we perform cross-validation with 10 folds to assess the model's performance across different subsets of the data.

```
rf_mod <- rand_forest(mode = "regression", trees = 1000) %>%
  set_engine("ranger")
rf_wf <-
  workflow() %>%
  add_model(rf_mod) %>%
  add_formula(Rating ~ .)
rf_fit <- rf_wf %>% fit(data = train_data)

folds <- vfold_cv(train_data, v = 10)

rf_fit_rs <-
  rf_wf %>%
  fit_resamples(folds)

rf_metrics <- collect_metrics(rf_fit_rs);rf_metrics
```

```
## # A tibble: 2 x 6
##   .metric .estimator mean      n std_err .config
##   <chr>   <chr>      <dbl> <int>  <dbl> <chr>
## 1 rmse    standard    0.208     10 0.00659 Preprocessor1_Model11
## 2 rsq     standard    0.533     10 0.0336  Preprocessor1_Model11
```

The cross-validation results for the random forest model demonstrate a RMSE of 0.193 with a standard error of 0.007 and an R-squared of 0.601, indicating strong predictive performance. For comparison, we specify and fit a linear regression model to the training data. Similar to the random forest model, cross-validation is performed to compare these models performance.

```
lm_recipe <- recipe(Rating ~ ., data = train_data)

lm_mod <- linear_reg() %>% set_engine("lm")

lm_wf <- workflow() %>%
  add_recipe(lm_recipe) %>%
  add_model(lm_mod)

lm_fit <- lm_wf %>% fit(data = train_data)

lm_fit_rs <- lm_wf %>% fit_resamples(folds)

lm_metrics <- collect_metrics(lm_fit_rs)
```

Here, we compare the performance metrics, between the random forest and linear regression models. This comparison provides insights into which model may be more suitable for predicting wine ratings.

```
comparison_metrics <- bind_rows(
  lm_metrics %>% mutate(Model = "Linear Regression"),
  rf_metrics %>% mutate(Model = "Random Forest"));comparison_metrics
```

```
## # A tibble: 4 x 7
```

##	.metric	.estimator	mean	n	std_err	.config	Model
##	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>	<chr>
## 1	rmse	standard	0.268	3	0.0110	Preprocessor1_Model11	Linear Regression
## 2	rsq	standard	0.235	3	0.0693	Preprocessor1_Model11	Linear Regression
## 3	rmse	standard	0.208	10	0.00659	Preprocessor1_Model11	Random Forest
## 4	rsq	standard	0.533	10	0.0336	Preprocessor1_Model11	Random Forest

The Random Forest model has a significantly lower RMSE and higher RSQ compared to the Linear Regression model, suggesting that it provides a more accurate prediction. \ Moreover, we further enhance the performance of our models by fine-tuning hyperparameters. Using grid search, we can fine-tune hyperparameters and select the best model configuration. By fine-tuning the Random Forest model, we get various combinations of hyperparameters.

```
tune_spec <- rand_forest(
  mtry = tune(),
  trees = 1000,
  min_n = tune()
) %>%
  set_mode("regression") %>%
  set_engine("ranger")

tree_wf <- workflow() %>%
  add_model(tune_spec) %>%
  add_formula(Rating ~ .)

tree_res <- tune_grid(
  tree_wf,
  resamples = folds,
  grid = 20
)
```

i Creating pre-processing data to finalize unknown parameter: mtry

```
mtry_range <- tree_res$.metrics[[1]] %>%
  select(mtry) %>%
  summarise(min_mtry = min(mtry),
            max_mtry = max(mtry))

min_n_range <- tree_res$.metrics[[1]] %>%
  select(min_n) %>%
  summarise(min_min_n = min(min_n),
            max_min_n = max(min_n))

rf_grid <- grid_regular(
  mtry(range = mtry_range),
  min_n(range = min_n_range),
  levels = 5
)

tree_res <- tune_grid(
  tree_wf,
  resamples = folds,
  grid = rf_grid
)

tree_res %>%
```

```
show_best(metric = "rmse")
```

```
## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean      n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     3    30 rmse     standard 0.207    10 0.00651 Preprocessor1_Model18
## 2     4    30 rmse     standard 0.207    10 0.00691 Preprocessor1_Model19
## 3     3    21 rmse     standard 0.207    10 0.00674 Preprocessor1_Model13
## 4     4    40 rmse     standard 0.207    10 0.00687 Preprocessor1_Model24
## 5     3    40 rmse     standard 0.207    10 0.00658 Preprocessor1_Model23
```

The `show_best()` function is used to identify the model with the lowest RMSE across different configurations. We then select the best model and perform the final fit using the optimized workflow, ensuring that the model is trained on the entire training dataset.

```
best_tree <- tree_res %>%
  select_best(metric = "rmse")
```

```
final_wf <-
  tree_wf %>%
  finalize_workflow(best_tree)
```

```
final_fit <-
  final_wf %>%
  last_fit(cell_split)
```

```
final_fit %>%
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>    <chr>      <dbl> <chr>
## 1 rmse     standard      0.198 Preprocessor1_Model1
## 2 rsq      standard      0.562 Preprocessor1_Model1
```

These are the metrics of our selected configuration, which should be performing reasonably in predicting wine rating.